

In [2]: `import numpy as np
import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator`

In [5]: `train_datagen = ImageDataGenerator(
 rescale=1./255,
 shear_range=0.2,
 zoom_range=0.2,
 horizontal_flip=True)
training_set = train_datagen.flow_from_directory(
 'training_set',
 target_size=(64, 64),
 batch_size=32,
 class_mode='categorical')`

Found 3459 images belonging to 5 classes.

In [6]: `test_datagen = ImageDataGenerator(rescale=1./255)
test_set = test_datagen.flow_from_directory(
 'test_set',
 target_size=(64, 64),
 batch_size=32,
 class_mode='categorical')`

Found 858 images belonging to 5 classes.

In [7]: `cnn = tf.keras.models.Sequential()`

In [8]: `cnn.add(tf.keras.layers.Conv2D(filters=64 , kernel_size=3 , activation='relu' , input_shape=[64,64,3]))
cnn.add(tf.keras.layers.MaxPool2D(pool_size=2, strides=2))`

C:\anaconda\Lib\site-packages\keras\src\layers\convolutional\base\_conv.py:107: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.  
super().\_\_init\_\_(activity\_regularizer=activity\_regularizer, \*\*kwargs)

In [9]: `cnn.add(tf.keras.layers.Conv2D(filters=64 , kernel_size=3 , activation='relu' ))
cnn.add(tf.keras.layers.MaxPool2D(pool_size=2 , strides=2))`

In [10]: `cnn.add(tf.keras.layers.Dropout(0.5))`

In [11]: `cnn.add(tf.keras.layers.Flatten())`

In [12]: `cnn.add(tf.keras.layers.Dense(units=128, activation='relu'))`

In [13]: `cnn.add(tf.keras.layers.Dense(units=5 , activation='softmax'))`

In [14]: `cnn.compile(optimizer = 'rmsprop' , loss = 'categorical_crossentropy' , metrics = ['accuracy'])`

In [15]: `cnn.fit(x = training_set , validation_data = test_set , epochs = 30)`

Epoch 1/30  
C:\anaconda\Lib\site-packages\keras\src\trainers\data\_adapters\py\_dataset\_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().\_\_init\_\_(\*\*kwargs)` in its constructor. `\*\*kwargs` can include `workers`, `use\_multiprocessing`, `max\_queue\_size`. Do not pass these arguments to `fit()`, as they will be ignored.  
self.\_warn\_if\_super\_not\_called()  
109/109 — 63s 552ms/step - accuracy: 0.3058 - loss: 1.5520 - val\_accuracy: 0.5361 - val\_loss: 1.0949  
Epoch 2/30  
109/109 — 16s 142ms/step - accuracy: 0.5196 - loss: 1.1580 - val\_accuracy: 0.6131 - val\_loss: 0.9763  
Epoch 3/30  
109/109 — 16s 141ms/step - accuracy: 0.5952 - loss: 1.0259 - val\_accuracy: 0.6212 - val\_loss: 0.9542  
Epoch 4/30  
109/109 — 16s 139ms/step - accuracy: 0.6158 - loss: 0.9672 - val\_accuracy: 0.6375 - val\_loss: 0.9167  
Epoch 5/30  
109/109 — 16s 143ms/step - accuracy: 0.6609 - loss: 0.8785 - val\_accuracy: 0.5862 - val\_loss: 1.1033  
Epoch 6/30  
109/109 — 17s 150ms/step - accuracy: 0.6665 - loss: 0.8694 - val\_accuracy: 0.6142 - val\_loss: 0.9769  
Epoch 7/30  
109/109 — 17s 146ms/step - accuracy: 0.7035 - loss: 0.8001 - val\_accuracy: 0.6200 - val\_loss: 1.0681  
Epoch 8/30  
109/109 — 16s 139ms/step - accuracy: 0.7059 - loss: 0.7737 - val\_accuracy: 0.6375 - val\_loss: 0.9107  
Epoch 9/30  
109/109 — 16s 144ms/step - accuracy: 0.7072 - loss: 0.7496 - val\_accuracy: 0.6037 - val\_loss: 1.1726  
Epoch 10/30  
109/109 — 22s 199ms/step - accuracy: 0.7232 - loss: 0.7204 - val\_accuracy: 0.6946 - val\_loss: 0.8329  
Epoch 11/30  
109/109 — 57s 509ms/step - accuracy: 0.7411 - loss: 0.6801 - val\_accuracy: 0.6970 - val\_loss: 0.8214  
Epoch 12/30  
109/109 — 75s 448ms/step - accuracy: 0.7469 - loss: 0.6450 - val\_accuracy: 0.6935 - val\_loss: 0.8562  
Epoch 13/30  
109/109 — 49s 432ms/step - accuracy: 0.7547 - loss: 0.6451 - val\_accuracy: 0.6620 - val\_loss: 0.9396  
Epoch 14/30  
109/109 — 51s 447ms/step - accuracy: 0.7866 - loss: 0.5850 - val\_accuracy: 0.6853 - val\_loss: 0.8712  
Epoch 15/30  
109/109 — 52s 462ms/step - accuracy: 0.7660 - loss: 0.6184 - val\_accuracy: 0.6702 - val\_loss: 0.8784  
Epoch 16/30  
109/109 — 52s 463ms/step - accuracy: 0.7883 - loss: 0.5676 - val\_accuracy: 0.6993 - val\_loss: 0.8656  
Epoch 17/30  
109/109 — 54s 486ms/step - accuracy: 0.7976 - loss: 0.5298 - val\_accuracy: 0.6503 - val\_loss: 0.9462  
Epoch 18/30  
109/109 — 53s 473ms/step - accuracy: 0.7996 - loss: 0.5105 - val\_accuracy: 0.6865 - val\_loss: 0.8822  
Epoch 19/30  
109/109 — 53s 475ms/step - accuracy: 0.8081 - loss: 0.5186 - val\_accuracy: 0.7249 - val\_loss: 0.8601  
Epoch 20/30  
109/109 — 60s 537ms/step - accuracy: 0.8093 - loss: 0.4854 - val\_accuracy: 0.6713 - val\_loss: 1.0125  
Epoch 21/30  
109/109 — 50s 446ms/step - accuracy: 0.8149 - loss: 0.4976 - val\_accuracy: 0.7168 - val\_loss: 0.8898  
Epoch 22/30  
109/109 — 54s 486ms/step - accuracy: 0.8317 - loss: 0.4575 - val\_accuracy: 0.6725 - val\_loss: 1.0386  
Epoch 23/30  
109/109 — 52s 460ms/step - accuracy: 0.8313 - loss: 0.4273 - val\_accuracy: 0.7040 - val\_loss: 0.9411  
Epoch 24/30  
109/109 — 52s 459ms/step - accuracy: 0.8443 - loss: 0.4176 - val\_accuracy: 0.6760 - val\_loss: 1.0296  
Epoch 25/30  
109/109 — 57s 512ms/step - accuracy: 0.8493 - loss: 0.4151 - val\_accuracy: 0.6678 - val\_loss: 1.0436  
Epoch 26/30  
109/109 — 53s 470ms/step - accuracy: 0.8552 - loss: 0.3910 - val\_accuracy: 0.6888 - val\_loss: 0.9671  
Epoch 27/30  
109/109 — 54s 475ms/step - accuracy: 0.8471 - loss: 0.4197 - val\_accuracy: 0.6970 - val\_loss: 0.9602  
Epoch 28/30  
109/109 — 68s 609ms/step - accuracy: 0.8607 - loss: 0.3625 - val\_accuracy: 0.7016 - val\_loss: 0.9873  
Epoch 29/30  
109/109 — 56s 499ms/step - accuracy: 0.8840 - loss: 0.3466 - val\_accuracy: 0.6888 - val\_loss: 1.1544  
Epoch 30/30  
109/109 — 57s 508ms/step - accuracy: 0.8637 - loss: 0.3555 - val\_accuracy: 0.7098 - val\_loss: 0.9345

Out[15]: <keras.src.callbacks.history.History at 0x15aa6b056d0>

In [34]: `from keras.preprocessing import image
test_image = image.load_img('Prediction/tu1.jpeg',target_size=(64,64))
test_image = image.img_to_array(test_image)
test_image = np.expand_dims(test_image,axis=0)
result = cnn.predict(test_image)`

1/1 — 0s 49ms/step

In [35]: `training_set.class_indices`

Out[35]: {'daisy': 0, 'dandelion': 1, 'rose': 2, 'sunflower': 3, 'tulip': 4}

In [36]: `if result[0][0]==1:
 print('Daisy')
elif result[0][1]==1:
 print('Dandelion')
elif result[0][2]==1:
 print('Rose')
elif result[0][3]==1:
 print('SunFlower')
elif result[0][4]==1:
 print("Tulip")`

Tulip

In [37]: `print(result)`

[[0. 0. 0. 0. 1.]]

In [ ]: